

Generative Models for Cat Image Synthesis

Deep Learning – Project III

Anna Ostrowska

June 2026

Abstract

This report presents the implementation and evaluation of three generative modelling approaches for synthesising cat images at 64×64 resolution: a Deep Convolutional GAN trained with Wasserstein loss and gradient penalty (DCGAN/WGAN-GP), a Vector-Quantized Variational Autoencoder with an autoregressive PixelCNN prior (VQ-VAE), and a Denoising Diffusion Probabilistic Model with DDIM fast sampling (DDPM). All models are trained from scratch on the CAT Dataset (9 997 images) using a single consumer GPU with 6 GiB VRAM. Models are compared quantitatively using the Fréchet Inception Distance (FID) computed on 5 000 generated images and qualitatively through visual inspection and latent space interpolation. DDPM achieves the lowest FID (24.36 at 50 DDIM steps with EMA weights), substantially outperforming VQ-VAE (187.54) and DCGAN (221.51). A post-hoc analysis of the DCGAN training configuration reveals that generator starvation caused by $n_{\text{critic}} = 5$ likely contributed to the high FID; corrected configurations reduce FID to 98.99 (standard DCGAN) and 158.48 (WGAN-GP $n_c = 2$). Ablation studies investigate the effect of DDIM inference steps and EMA weights on DDPM quality, and the sensitivity of DCGAN to latent dimension and learning rate. An exploratory extension trains DCGAN on a combined cats-and-dogs dataset, finding a suggestive improvement in cat FID that warrants further investigation. Source code and model checkpoints are provided for full reproducibility.

Contents

1	Introduction	3
2	Background and Related Work	3
2.1	Generative Adversarial Networks	3
2.2	Variational Autoencoders and VQ-VAE	3
2.3	Denoising Diffusion Probabilistic Models	4
2.4	Fréchet Inception Distance	4
3	Dataset and Preprocessing	4
3.1	Cat Dataset	4
3.2	Preprocessing Pipeline	4
3.3	Data Augmentation	5

4	Methods	5
4.1	DCGAN with WGAN-GP	5
4.2	VQ-VAE + PixelCNN Prior	6
4.3	DDPM with DDIM Sampling	6
5	Computational Considerations	7
6	Results	7
6.1	FID Comparison	7
6.2	Generated Samples	8
6.3	Training Dynamics	8
6.4	Sample Progression	9
6.5	VQ-VAE Reconstruction	10
6.6	Latent Interpolation	10
7	Mode Collapse Analysis	11
8	DCGAN Training Bottleneck Analysis	12
8.1	Generator Starvation	12
8.2	Corrected Configurations	12
8.3	Discussion	13
9	Ablation Studies	13
9.1	DDPM: DDIM Inference Steps	13
9.2	DDPM: EMA Weights at Inference	14
9.3	DCGAN: Latent Dimension and Learning Rate	14
10	Cats and Dogs Experiment	15
10.1	Setup	15
10.2	Evaluation	15
10.3	Discussion	16
11	Conclusion	17

1 Introduction

This report presents the design, training, and evaluation of three classes of generative image models applied to a dataset of cat photographs. The three approaches investigated are:

1. **DCGAN** with Wasserstein loss and gradient penalty (WGAN-GP),
2. **VQ-VAE** with an autoregressive PixelCNN prior,
3. **DDPM** (Denoising Diffusion Probabilistic Model) with DDIM sampling.

Models are compared quantitatively using the Fréchet Inception Distance (FID) and qualitatively through visual inspection of generated samples and latent interpolation strips. Additionally, an ablation study is conducted for each model class, and a supplementary experiment trains a DCGAN on a combined cats-and-dogs dataset.

2 Background and Related Work

2.1 Generative Adversarial Networks

Generative Adversarial Networks (GANs) [3] frame image synthesis as a two-player game between a Generator G and a Discriminator D . The generator maps samples $\mathbf{z} \sim p(\mathbf{z})$ from a simple prior (typically $\mathcal{N}(0, I)$) to synthetic images, while the discriminator attempts to distinguish real from generated data. Training minimises the standard min-max objective:

$$\min_G \max_D \mathbb{E}_{x \sim p_r} [\log D(x)] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))] \quad (1)$$

GANs are prone to *mode collapse* – a failure mode where the generator produces only a small subset of the data distribution – and to unstable gradient dynamics. The Deep Convolutional GAN (DCGAN) [10] introduced architectural guidelines (all-convolutional networks, batch normalisation, LeakyReLU activations) that substantially improved training stability.

WGAN and WGAN-GP. Arjovsky et al. [1] proposed replacing the Jensen-Shannon divergence objective with the Wasserstein-1 distance, which provides smoother gradients and alleviates mode collapse. However, their weight-clipping approach introduces gradient pathologies. Gulrajani et al. [4] resolved this by enforcing the Lipschitz constraint via a gradient penalty term on interpolated samples, yielding the Wasserstein GAN with Gradient Penalty (WGAN-GP). Crucially, the WGAN-GP critic outputs an unbounded scalar score rather than a probability, so binary label targets such as $\{0, 1\}$ or their smoothed variants are not used in the Wasserstein loss path.

2.2 Variational Autoencoders and VQ-VAE

Variational Autoencoders (VAEs) [9] train an encoder–decoder pair by maximising the evidence lower bound (ELBO):

$$\mathcal{L}_{\text{VAE}} = \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \| p(z)), \quad (2)$$

where q_ϕ is the approximate posterior and p_θ the decoder. The continuous latent space allows smooth interpolation but tends to produce blurry reconstructions.

The Vector-Quantized VAE (VQ-VAE) [13] replaces the continuous latent code with a discrete codebook. The encoder output is replaced by the nearest codebook entry (straight-through gradient); codebook updates use exponential moving averages (EMA) for stability. Because the prior $p(z)$ is non-trivial (discrete, structured), a separate autoregressive model (PixelCNN [12]) is trained over the quantised token sequences after the autoencoder converges.

2.3 Denoising Diffusion Probabilistic Models

DDPM [6] defines a *forward process* that gradually adds Gaussian noise to data:

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I), \quad (3)$$

and a learnable *reverse process* $p_\theta(x_{t-1}|x_t)$ implemented as a U-Net that predicts the noise ϵ added at each step. The training objective is:

$$\mathcal{L}_{\text{DDPM}} = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(x_t, t)\|^2]. \quad (4)$$

Dhariwal and Nichol [2] demonstrated that diffusion models with classifier guidance surpass GANs on FID benchmarks.

DDIM. Song et al. [11] showed that the same trained model can be used with a deterministic, non-Markovian sampling process (Denoising Diffusion Implicit Models) requiring as few as 20–50 function evaluations rather than the full $T = 1000$ steps, at negligible quality loss.

2.4 Fréchet Inception Distance

FID [5] measures the Fréchet (Wasserstein-2) distance between the distributions of real and generated images in the feature space of an Inception-v3 network:

$$\text{FID} = \|\mu_r - \mu_g\|^2 + \text{tr}(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2}), \quad (5)$$

where (μ_r, Σ_r) and (μ_g, Σ_g) are the mean and covariance of the Inception activations for real and generated sets. Lower FID indicates closer distributional similarity. We use `torch-fidelity` with 5 000 real and 5 000 generated images.

3 Dataset and Preprocessing

3.1 Cat Dataset

The primary dataset is the CAT Dataset [14], consisting of 9 997 RGB cat photographs with heterogeneous resolutions, organised in subdirectories `CAT_00` through `CAT_06`.

3.2 Preprocessing Pipeline

A uniform preprocessing pipeline was applied to all images:

1. **Resize** the shorter edge to 64 pixels.
2. **CenterCrop** to 64×64 .

3. **Normalize** to $[-1, 1]$ via $x \leftarrow (x - 0.5)/0.5$ per channel.

All models operate on $64 \times 64 \times 3$ tensors. Figure 1 shows a random selection of preprocessed images. A fixed random seed (42) is used for all experiments. DataLoaders use `drop_last=True`, giving $\lfloor 9997/64 \rfloor = 156$ batches per epoch.



Figure 1: Random sample of real cat images after 64×64 center-crop preprocessing.

3.3 Data Augmentation

An optional augmentation pipeline was implemented. When enabled, the following transforms are applied before normalization:

- `RandomResizedCrop` (scale 0.85–1.0),
- `RandomHorizontalFlip` ($p = 0.5$),
- `ColorJitter` (brightness/contrast/saturation = 0.2, hue = 0.05),
- `RandomGrayscale` ($p = 0.05$),
- `RandomRotation` ($\pm 10^\circ$),
- `RandomErasing` ($p = 0.3$, scale 2–8%).

The baseline models are trained without augmentation; augmentation effects are investigated in the ablation study (Section 9).

4 Methods

4.1 DCGAN with WGAN-GP

Architecture. The Generator G maps a latent vector $\mathbf{z} \sim \mathcal{N}(0, I)$, $\mathbf{z} \in \mathbb{R}^{128}$, to a $64 \times 64 \times 3$ image through four transposed-convolution blocks (BatchNorm + ReLU), ending in a Tanh activation. The Discriminator D mirrors this with four strided-convolution blocks (LeakyReLU, $\alpha = 0.2$) and outputs a scalar logit without sigmoid. The architecture follows DCGAN guidelines [10].

Training objective. We adopt the Wasserstein loss with gradient penalty (WGAN-GP [4]):

$$\mathcal{L}_D = \mathbb{E}_{\tilde{x} \sim P_g}[D(\tilde{x})] - \mathbb{E}_{x \sim P_r}[D(x)] + \lambda \mathbb{E}_{\hat{x}}[(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2], \quad (6)$$

with $\lambda = 10$ and $\hat{x}$ sampled uniformly along straight lines between real and generated data. The discriminator (critic) is updated $n_{\text{critic}} = 5$ times per generator step. Because the WGAN-GP critic returns an unbounded scalar rather than a class probability, binary label targets – and any smoothing thereof – have no role in this loss formulation and were not applied.

Hyperparameters. Trained for 100 epochs, batch size 64, Adam ($\beta_1 = 0.5$, $\beta_2 = 0.999$), learning rate 2×10^{-4} for both G and D .

4.2 VQ-VAE + PixelCNN Prior

Architecture. The VQ-VAE [13] consists of an encoder E that maps a 64×64 image to a 16×16 discrete codebook index map, a vector-quantization layer with codebook size $K = 512$ and embedding dimension $d = 64$, and a decoder D that reconstructs the image from quantized embeddings. The encoder and decoder each use strided-convolution and transposed-convolution blocks with two residual layers. Codebook updates use exponential moving averages (EMA) with decay $\gamma = 0.99$.

Two-phase training:

1. **Phase 1 – Autoencoder** (100 epochs): minimizes reconstruction loss + VQ commitment loss $\beta \|z_e - \text{sg}(z_q)\|^2$, $\beta = 0.25$.
2. **Phase 2 – Prior** (100 epochs): a PixelCNN [12] with 15 masked residual layers (128 filters each) is trained autoregressively on the flattened $16 \times 16 = 256$ token sequences from the frozen encoder.

Sampling. At inference, the PixelCNN samples a 256-token sequence autoregressively; tokens are mapped to embeddings via the codebook and decoded by D .

4.3 DDPM with DDIM Sampling

Architecture. We implement the DDPM framework [6] with a U-Net backbone. The U-Net has channel multipliers $(1, 2, 2, 2)$, two residual blocks per resolution, and self-attention at the 16×16 feature map level. An exponential moving average (EMA, decay 0.9999) of model weights is maintained and used at inference.

Diffusion schedule. Linear noise schedule with $T = 1000$ steps, $\beta_1 = 10^{-4}$, $\beta_T = 0.02$. The forward process is:

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I), \quad (7)$$

where $\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$. The model predicts the added noise $\epsilon_\theta(x_t, t)$ and is trained with simple ℓ_2 loss.

DDIM Sampling. Deterministic DDIM [11] allows generating samples with far fewer function evaluations than the full $T = 1000$ reverse chain. At inference we use 50 DDIM steps ($\eta = 0.0$) by default, giving a $20\times$ speedup.

Hyperparameters. Trained for 200 epochs, batch size 16 (limited by 6 GiB VRAM), Adam ($\beta_1 = 0.9$, $\beta_2 = 0.999$), learning rate 2×10^{-4} with linear warmup.

5 Computational Considerations

All experiments were run on a single NVIDIA GPU with 6 GiB VRAM and a Windows host. Key adaptations:

- `num_workers=0` in all DataLoaders (Windows multiprocessing limitation).
- DDPM batch size reduced to 16 to fit within VRAM.
- DDIM sampling used instead of full reverse diffusion to cut inference time from ≈ 8 h (1000 steps) to ≈ 25 min (50 steps).
- VQ-VAE EMA codebook uses `scatter_add_ + bincount` instead of `F.one_hot` to avoid out-of-memory issues.
- Global seed 42 for reproducibility.

6 Results

6.1 FID Comparison

Fréchet Inception Distance (FID) [5] measures the distance between the distribution of real images and generated images in the feature space of an Inception-v3 network. We generate 5000 samples per model and compare against 5000 real images drawn from the training set.

Table 1: FID scores on the cat dataset (64×64 , 5000 samples).

Model	FID ↓	Notes
DCGAN (WGAN-GP)	221.51	100 epochs, single run
VQ-VAE + PixelCNN	187.54	100 + 100 epochs, single run
DDPM (50 DDIM steps)	24.36	200 epochs, EMA, single run

DDPM achieves a substantially lower FID, consistent with the literature showing diffusion models outperforming GANs and VAEs on small-scale datasets [2].

Scope of this comparison. Two caveats apply. First, the comparison is not compute-normalised: DDPM was trained for ≈ 12 h, DCGAN for ≈ 2 h, and VQ-VAE for ≈ 4 h. The table reflects the best feasible configuration for each model under the 6 GiB VRAM constraint, not a strictly equal-budget comparison. Second, FID was computed against a subset of the *training* distribution; no held-out test set was used. The metric therefore measures similarity to available training data rather than generalisation to unseen cat images. Given that the main comparisons are single-run, the large gap between DDPM and the other models is reliably qualitative, while smaller differences should not be over-interpreted.

6.2 Generated Samples




Figure 2: Generated samples from each model at the end of training.

DDPM samples show the most coherent cat structure with natural fur textures. DCGAN samples exhibit some diversity but with noticeable artifacts and occasional mode collapse symptoms (similar face structures). VQ-VAE samples are globally plausible but blurrier, which likely reflects limitations of the PixelCNN prior rather than the autoencoder itself (see Section 6.5).

6.3 Training Dynamics




Figure 3: Training loss curves for all three models.




Figure 4: VQ-VAE PixelCNN prior training loss (Phase 2, 100 epochs).

DCGAN training shows the characteristic adversarial oscillation: generator and discriminator losses compete without a stable equilibrium. VQ-VAE reconstruction loss decreases steadily throughout Phase 1 and the PixelCNN prior loss decreases more slowly in Phase 2. DDPM loss decreases monotonically and smoothly, reflecting the well-defined mean-squared-error noise-prediction objective.

All models were trained on the full dataset without a validation split; early stopping was therefore not applied. Sample grids saved at regular checkpoints (shown in Section 6.4) serve as the primary qualitative proxy for monitoring training progress and detecting mode collapse. Per-epoch FID evaluation was not feasible due to its computational cost (≈ 5 min per evaluation).

6.4 Sample Progression




Figure 5: DCGAN sample grids at training epochs 10, 50, and 100.




Figure 6: DDPM sample grids at training epochs 10, 100, and 200.

6.5 VQ-VAE Reconstruction



Figure 7: VQ-VAE reconstructions at epoch 100 (top row: inputs, bottom row: reconstructed).

The VQ-VAE achieves a codebook perplexity of 403.3/512, indicating that $\approx 79\%$ of codes are actively used. Reconstruction loss at convergence is 0.0066 (MSE), VQ commitment loss 0.0033. The good reconstruction quality contrasts with the blurrier unconditional samples, which suggests that the PixelCNN prior is a likely bottleneck: the autoencoder itself captures sufficient detail, but the prior’s ability to model the full distribution of token sequences is limited.


6.6 Latent Interpolation

To assess the smoothness of the learned latent space, we interpolate between two latent vectors $\mathbf{z}_A$ and $\mathbf{z}_B$. For models whose prior is $\mathcal{N}(0, I)$ (DCGAN, DDPM), we use *spherical*

linear interpolation (slerp):

$$\mathbf{z}_\alpha = \frac{\sin((1-\alpha)\theta)}{\sin\theta} \mathbf{z}_A + \frac{\sin(\alpha\theta)}{\sin\theta} \mathbf{z}_B, \quad \theta = \arccos\left(\frac{\mathbf{z}_A \cdot \mathbf{z}_B}{\|\mathbf{z}_A\| \|\mathbf{z}_B\|}\right), \quad (8)$$


for $\alpha \in \{0, \frac{1}{9}, \dots, 1\}$. Slerp preserves the norm of vectors drawn from an isotropic Gaussian: linear interpolation through the origin would traverse a low-density region of $\mathcal{N}(0, I)$ and produce systematically blurry intermediate frames. For VQ-VAE, we interpolate linearly in the continuous pre-quantization encoder space, since the prior operates over discrete token indices rather than a Gaussian.



(a) DCGAN: spherical interpolation (slerp) in $\mathcal{N}(0, I)$ latent space.



(b) VQ-VAE: linear interpolation in continuous pre-quantization encoder space.



(c) DDPM: spherical interpolation (slerp) between two initial noise tensors $x_T \sim \mathcal{N}(0, I)$.

Figure 8: Latent interpolation strips (10 images: 2 endpoints + 8 intermediate frames).

Discussion. Smooth interpolations indicate that the model has learned a continuous latent manifold rather than memorising isolated training examples. DCGAN shows relatively smooth slerp transitions; intermediate images are recognisable as cat faces. VQ-VAE interpolation is smooth in encoder space but the discrete quantization step introduces occasional abrupt transitions when token assignments change near the boundary of two Voronoi cells. DDPM slerp in x_T space produces the most photorealistic intermediate frames, which is expected since the diffusion reverse process is well-conditioned on the initial noise tensor.

7 Mode Collapse Analysis

DCGAN’s high FID (221.51) is partly explained by mode collapse. We observe:

- Many generated images show similar facial proportions and color palettes, suggesting the generator maps large regions of $\mathcal{N}(0, I)$ to similar outputs.
- Loss curves show the characteristic WGAN-GP pattern: the critic loss stabilises, but the generator loss remains volatile.

- Despite WGAN-GP mitigating collapse compared to vanilla GAN, the diversity of generated images remains lower than real data.

Mitigation strategies applied:

- WGAN-GP gradient penalty ($\lambda = 10$) replaces standard discriminator weight clipping.
- $n_{\text{critic}} = 5$ discriminator steps per generator step, keeping the critic well-calibrated.
- Note: the shared training code contained a label smoothing parameter (real labels $\rightarrow 0.9$) intended for standard BCE-GAN training. This parameter operates on binary cross-entropy targets and has no effect on the WGAN-GP loss path, which uses a scalar critic score rather than binary labels. It was therefore *inactive* during WGAN-GP training.

The root cause of the high FID is analysed in more detail in Section 8.

8 DCGAN Training Bottleneck Analysis

8.1 Generator Starvation

A post-hoc analysis of the DCGAN training budget revealed a potential source of under-training. With $n_{\text{critic}} = 5$ and 156 batches per epoch ($\lfloor 9997/64 \rfloor$ with `drop_last=True`), the generator receives only $\lfloor 156/5 \rfloor \times 100 = \mathbf{3\,100}$ update steps over 100 epochs. A standard DCGAN with $n_c = 1$ would accumulate $156 \times 100 = \mathbf{15\,600}$ generator updates at the same computational cost.

8.2 Corrected Configurations

Two corrected configurations were trained to test whether increasing the generator update budget improves FID:

1. **Standard DCGAN**: switch to standard BCE loss with $n_c = 1$ (the canonical DCGAN setting, not a compromise); increased generator capacity ($g_{\text{ch}} = 128$); 150 epochs $\Rightarrow \mathbf{23\,400}$ G-steps.
2. **WGAN-GP** $n_c = 2$: pragmatic reduction of critic steps while retaining the Wasserstein loss; 200 epochs $\Rightarrow \mathbf{15\,600}$ G-steps. Note that $n_c < 5$ weakens the Lipschitz guarantee; a theoretically proper fix would require ≈ 500 epochs.

Table 2: FID results for original and corrected DCGAN configurations.

Configuration	G-steps	FID ↓	Notes
Original WGAN-GP ($n_c = 5$, 100 ep.)	3 100	221.51	baseline
Standard DCGAN ($n_c = 1$, 150 ep.)	23 400	98.99	BCE loss
WGAN-GP $n_c = 2$ (200 ep.)	15 600	158.48	Wasserstein loss




Figure 9: Samples from corrected DCGAN configurations.

8.3 Discussion

Both corrected configurations yield substantially lower FID. The standard DCGAN at 23 400 G-steps achieves FID 98.99, a $2.2\times$ improvement over the original. The WGAN-GP $n_c = 2$ at 15 600 G-steps achieves FID 158.48.

These results suggest that the original WGAN-GP result was limited not only by architecture but also by the low training budget allocated to generator updates. However, because each corrected configuration changes several factors simultaneously – loss function, n_c , model capacity, and training duration – these should be interpreted as diagnostic follow-up experiments rather than a clean single-factor ablation. The improvement from standard DCGAN over WGAN-GP $n_c = 2$ may reflect the larger generator update budget (23 400 vs. 15 600 steps), the switch to BCE loss, or increased generator capacity, and cannot be attributed to any single factor from this experiment alone.

9 Ablation Studies

9.1 DDPM: DDIM Inference Steps

We evaluate the effect of the number of DDIM steps on FID without retraining (inference-only ablation on the epoch-200 checkpoint, non-EMA weights).

Table 3: DDPM FID vs. number of DDIM inference steps (epoch-200 checkpoint, non-EMA weights, 5000 samples).

DDIM Steps	FID ↓	Speedup vs. 1000-step
10	66.92	100×
25	60.02	40×
50	54.51	20×
100	48.70	10×
200	39.81	5×

FID improves monotonically with more steps, with diminishing returns above 100. The 50-step setting (20× speedup, FID 54.5) provides a practical speed-quality trade-off. Steps below 25 degrade noticeably (+27 FID vs. 50 steps), making them unsuitable for qualitative evaluation.

9.2 DDPM: EMA Weights at Inference

Using EMA weights at inference has a large effect independent of step count.

Table 4: Effect of EMA weights at 50 DDIM steps (epoch-200 checkpoint).

Inference weights	FID ↓
Without EMA	54.51
With EMA (decay 0.9999)	24.36

EMA weights reduce FID by 55% at 50 steps (from 54.51 to 24.36), a larger improvement than doubling the number of inference steps. The EMA averages weight noise that accumulates during training, which is particularly beneficial for diffusion models because errors in the reverse chain compound over many steps. The 50-step EMA configuration is therefore the default reported throughout this work.

9.3 DCGAN: Latent Dimension and Learning Rate

We train DCGAN configurations for 100 epochs each, varying the latent dimension z_{dim} and learning rate independently across three seeds per configuration (all other hyperparameters fixed to baseline: $z_{\text{dim}} = 128$, $\text{lr} = 2 \times 10^{-4}$, $n_c = 5$). FID is reported as mean $\pm$ population standard deviation across seeds.

Table 5: DCGAN ablation: FID (mean $\pm$ std across 3 seeds, 100 epochs).

Configuration	FID mean $\pm$ std	Δ vs. baseline
$z_{\text{dim}} = 64$	246.65 ± 5.57	-17.60
$z_{\text{dim}} = 128$ (baseline)	264.25 ± 20.21	$-$
$z_{\text{dim}} = 256$	266.05 ± 2.52	$+1.80$
$\text{lr} = 1 \times 10^{-4}$	295.11 ± 18.27	$+30.86$
$\text{lr} = 2 \times 10^{-4}$ (baseline)	264.25 ± 20.21	$-$
$\text{lr} = 4 \times 10^{-4}$	220.05 ± 4.47	-44.20
Augmented (z_{128} , $\text{lr } 2 \times 10^{-4}$)	284.05 ± 14.80	$+19.80$

Latent dimension has a small effect: the differences between $z_{\text{dim}} \in \{64, 128, 256\}$ are within the baseline seed variance (± 20), so no firm conclusion about the optimal z_{dim} can be drawn. The high baseline variance itself is notable and indicates that DCGAN training outcomes are sensitive to random initialisation.

Learning rate has a much larger impact. A higher lr of 4×10^{-4} reduces mean FID by 44 points with low variance (± 4.47), while the lowest lr (1×10^{-4}) increases mean FID by 31 points. The likely mechanism is that a low learning rate causes the generator to update too slowly to keep pace with the well-trained critic, resulting in poor gradient signal.

Enabling data augmentation hurts FID (284.05 ± 14.80 vs. 264.25 ± 20.21 baseline). This is consistent with observations in the GAN literature: when the discriminator is trained on augmented (distorted) real images without applying the same augmentations to generated images, it learns to accept distortions as “real”, which incorrectly guides the generator toward augmentation artefacts rather than the true data distribution. Discriminator-adaptive augmentation strategies such as ADA [8] are designed to address this problem but were not implemented here.

10 Cats and Dogs Experiment

10.1 Setup


The Kaggle dogs-vs-cats dataset [7] provides 12 500 cat images and 12 500 dog images (25 000 total). We train the same DCGAN architecture (100 epochs, all hyperparameters identical to the cats-only baseline) on the combined dataset without class labels, as an unconditional generative model.

10.2 Evaluation

FID is computed separately against 5 000 real cat images and 5 000 real dog images to measure how well the generated distribution covers each class.

Table 6: DCGAN (cats+dogs) FID compared to cats-only model.

Model	FID vs. Cats ↓	FID vs. Dogs ↓	FID vs. Combined ↓
DCGAN (cats only, 100 ep.)	221.51	–	–
DCGAN (cats+dogs, 100 ep.)	206.57	197.79	190.39



(a) Epoch 25

(b) Epoch 100

Figure 10: DCGAN (cats+dogs) sample grids at epochs 25 and 100.

10.3 Discussion

Training an unconditional generative model on a two-class dataset poses an interesting challenge: without conditioning, the model cannot guarantee class-distinct generation.

The cats+dogs DCGAN achieves FID 206.57 against real cats, 15 points lower than the cats-only model’s FID of 221.51. This improvement is suggestive of a beneficial effect: additional dog images may expose the model to a wider variety of animal textures and structures, with shared low-level features (fur, eyes, snout shape) acting as implicit data augmentation. However, this interpretation should be treated cautiously. The 15-point improvement is comparable to the seed variance observed in the DCGAN ablation ($\sigma \approx 20$ FID), so the result is not statistically conclusive from a single run. Furthermore, FID alone cannot confirm whether the generated images are visually distinct as cats or dogs versus blended representations; a classifier-based analysis would be needed to verify class identity in generated samples.

The FID against real dogs (197.79) is similar to the cat FID (206.57), suggesting the model does not strongly collapse to one class.

As a more direct measure of how well the generated distribution covers the training distribution, we also computed FID against a combined reference set of 2500 real cats and 2500 real dogs ($N = 5000$ total). The resulting FID is 190.39, lower than either individual class FID. This is the most methodologically appropriate comparison for an

unconditional model trained on a mixed dataset: the reference set reflects the same class balance as the training data. The improvement over the per-class FIDs (from ≈ 202 average to 190) is consistent with the model covering a blended distribution rather than cleanly separating into cat and dog modes.

11 Conclusion

We trained and evaluated three families of generative models on cat images:

- **DDPM** achieves the best FID (24.36 with EMA at 50 DDIM steps) and the most visually coherent samples. The 50-step DDIM configuration provides a practical speed-quality trade-off ($20\times$ speedup over full sampling). EMA weights are the single most impactful component: they reduce FID by 55% relative to non-EMA inference at the same step count.
- **VQ-VAE + PixelCNN** achieves FID 187.54. Reconstruction quality is good (MSE = 0.0066, codebook utilisation 79%), but the gap between reconstruction quality and sample quality suggests the PixelCNN prior is a likely bottleneck for unconditional generation.
- **DCGAN (WGAN-GP)** achieves FID 221.51, with visible mode collapse symptoms despite WGAN-GP regularisation. Training is the fastest (~ 2 h). Post-hoc analysis suggests that generator starvation ($n_c = 5$ yielding only 3 100 G-steps) contributed to the poor result; corrected configurations reduce FID to 98.99 (standard DCGAN, 23 400 G-steps) and 158.48 (WGAN-GP $n_c = 2$, 15 600 G-steps).

These results are consistent with the literature showing that diffusion models offer superior image quality compared to GAN and VAE approaches on small-scale image datasets with limited compute [2].

Methodological limitations. The main FID comparisons are single-run results; the relative ranking is reliable for large gaps but the VQ-VAE vs. DCGAN difference (≈ 34 FID) may not be robust. The comparison is not compute-normalised. No held-out test set was used: FID reflects similarity to the training distribution. The cats-and-dogs improvement and the augmentation penalty are both within the observed seed variance ($\approx \pm 20$ FID) and should therefore be treated as exploratory findings. The corrected DCGAN experiments change multiple factors simultaneously, so the improvement cannot be attributed to generator starvation alone.

Potential improvements and future work.

- Training DDPM at 128×128 resolution with gradient checkpointing would likely improve FID further.
- A class-conditional DCGAN (e.g., projection discriminator) trained on the cats-and-dogs dataset would enable targeted cat/dog generation and allow a cleaner assessment of multi-class transfer.
- Replacing the PixelCNN prior with a Transformer-based prior (e.g., VQ-VAE-2 [13]) could improve sample diversity.
- A cosine noise schedule for DDPM may improve results at lower resolutions.

AI Usage Statement

GitHub Copilot (Claude Sonnet) was used as a coding assistant throughout this project for: boilerplate generation, debugging PyTorch errors, and suggestions for hyperparameter configurations. All architectural decisions, experimental design, and analysis were performed by the author. Generated code was reviewed and adapted before use.

Reproducibility

All code is available in the project repository. To reproduce results:

1. Install dependencies: `pip install -r requirements.txt`
2. Download and place the CAT Dataset in `data/cats/`
3. Train models: `python -m src.training.train_dcgan -config configs/dcgan_config.yaml`, analogously for VQ-VAE and DDPM.
4. Evaluate: `python evaluate_all.py`
5. Run ablations: `python ablate_ddpm_steps.py` and `python ablate_dcgan_multiseed.py`

Global seed 42 is set via `src/utils/seed.py` for full reproducibility.

References

- [1] Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein generative adversarial networks. In *International Conference on Machine Learning*, pages 214–223, 2017.
- [2] Prafulla Dhariwal and Alexander Nichol. Diffusion models beat GANs on image synthesis. *Advances in Neural Information Processing Systems*, 34:8780–8794, 2021.
- [3] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in Neural Information Processing Systems*, volume 27, 2014.
- [4] Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, and Aaron Courville. Improved training of wasserstein GANs. *Advances in Neural Information Processing Systems*, 30, 2017.
- [5] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. GANs trained by a two time-scale update rule converge to a local Nash equilibrium. *Advances in Neural Information Processing Systems*, 30, 2017.
- [6] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems*, 33:6840–6851, 2020.
- [7] Kaggle. Dogs vs. cats, 2013. Available at <https://www.kaggle.com/competitions/dogs-vs-cats/>.
- [8] Tero Karras, Miika Aittala, Janne Hellsten, Samuli Laine, Jaakko Lehtinen, and Timo Aila. Training generative adversarial networks with limited data. In *Advances in Neural Information Processing Systems*, volume 33, pages 12104–12114, 2020.

- [9] Diederik P. Kingma and Max Welling. Auto-encoding variational Bayes. *International Conference on Learning Representations*, 2014.
- [10] Alec Radford, Luke Metz, and Soumith Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. *International Conference on Learning Representations*, 2016.
- [11] Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. *International Conference on Learning Representations*, 2021.
- [12] Aaron van den Oord, Nal Kalchbrenner, and Koray Kavukcuoglu. Pixel recurrent neural networks. In *International Conference on Machine Learning*, pages 1747–1756, 2016.
- [13] Aaron van den Oord, Oriol Vinyals, and Koray Kavukcuoglu. Neural discrete representation learning. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- [14] Weiwei Zhang, Jian Sun, and Xiaoou Tang. Cat head detection – CAT dataset, 2008. Available at <https://www.kaggle.com/datasets/crawford/cat-dataset>.